

Big Data Project

Wikimedia recent change streaming parsing

Khassangali Gapparov 2022 @MIU

Technology stack

Overall main technologies list

- Streaming source from Wikimedia recent change: <https://stream.wikimedia.org/v2/stream/recentchange>
- Message queue tool - Kafka (local cluster)
- Streaming source event listener okHttp3.event source
- Spark with Streaming API
- HDFS cluster to save the final output
- IntelliJ Idea for creating java applications (part1 consumer, part2 producer)

Part1

Kafka Producer Application

- Created scala project with appropriate libraries:

```
ThisBuild / version := "0.1.0-SNAPSHOT"

ThisBuild / scalaVersion := "2.11.12"

lazy val root = (project in file("."))
  .settings(
    libraryDependencies ++= Seq(
      "org.apache.spark" %% "spark-streaming" % "2.4.8",
      "org.apache.spark" %% "spark-core" % "2.4.8",
      // https://mvnrepository.com/artifact/org.apache.kafka/kafka-clients
    )
    libraryDependencies += "org.apache.kafka" % "kafka-clients" % "3.2.1"
    // https://mvnrepository.com/artifact/org.apache.spark/spark-streaming-kafka-0-10
    libraryDependencies += "org.apache.spark" %% "spark-streaming-kafka-0-10" % "2.4.8"
```

Part 1 cont...

Http event source

- Listen the URL for events and @onMessage we send message to Kafka producer with content of payload.

```
@Configuration
@PropertySource("classpath:kafka.properties")
public class WikiEventSource {

    private final WikiMediaHandler wikiMediaHandler;

    @Value("https://stream.wikimedia.org/v2/stream/recentchange")
    private String url;

    @Autowired
    public WikiEventSource(WikiMediaHandler wikiMediaHandler) { this.wikiMediaHandler = wikiMediaHandler; }

    @Bean
    public EventSource eventSource() { return new EventSource.Builder(wikiMediaHandler, URI.create(url)).build(); }
```

```
@Component
public class WikiMediaHandler implements EventHandler {

    public static final Logger log = LoggerFactory.getLogger(WikiMediaHandler.class.getSimpleName());

    private final KafkaTemplate<String, String> kafkaTemplate;

    private final NewTopic topic;

    @Autowired
    public WikiMediaHandler(KafkaTemplate<String, String> producer, NewTopic topic) {...}

    @Override
    public void onOpen() {}

    @Override
    public void onClosed() {}

    @Override
    public void onMessage(String event, MessageEvent messageEvent) throws Exception {
        kafkaTemplate.send(topic.name(), messageEvent.getData());
    }
}
```

unsmoker, 9/12/22, 3:20 PM • KafkaProduc

Part 2

Kafka Consumer

Part2

Kafka Consumer

- We subscribe to the Kafka topic (“Wikimedia.recentchange”). with Kafka params shown as bellow configuration:

```
public class StreamingContext implements Serializable {

    private static final JavaStreamingContext javaStreamingContext = new JavaStreamingContext(new SparkConf().setMaster("local[2]").setAppName("SparkStreaming"), Durations.seconds(2));

    private static final Map<String, Object> kafkaParams = new HashMap<>();

    public static JavaInputDStream<ConsumerRecord<String, String>> javaInputDStream(){
        kafkaParams.put("bootstrap.servers", "localhost:9092");
        kafkaParams.put("key.deserializer", StringDeserializer.class);
        kafkaParams.put("value.deserializer", StringDeserializer.class);
        kafkaParams.put("group.id", "spark_consumer_group");

        return KafkaUtils.createDirectStream(
            javaStreamingContext,
            LocationStrategies.PreferConsistent(),
            ConsumerStrategies.Subscribe(Collections.singleton("wikimedia.recentchange"), kafkaParams)
        );
    }
}
```


Kafka consumer

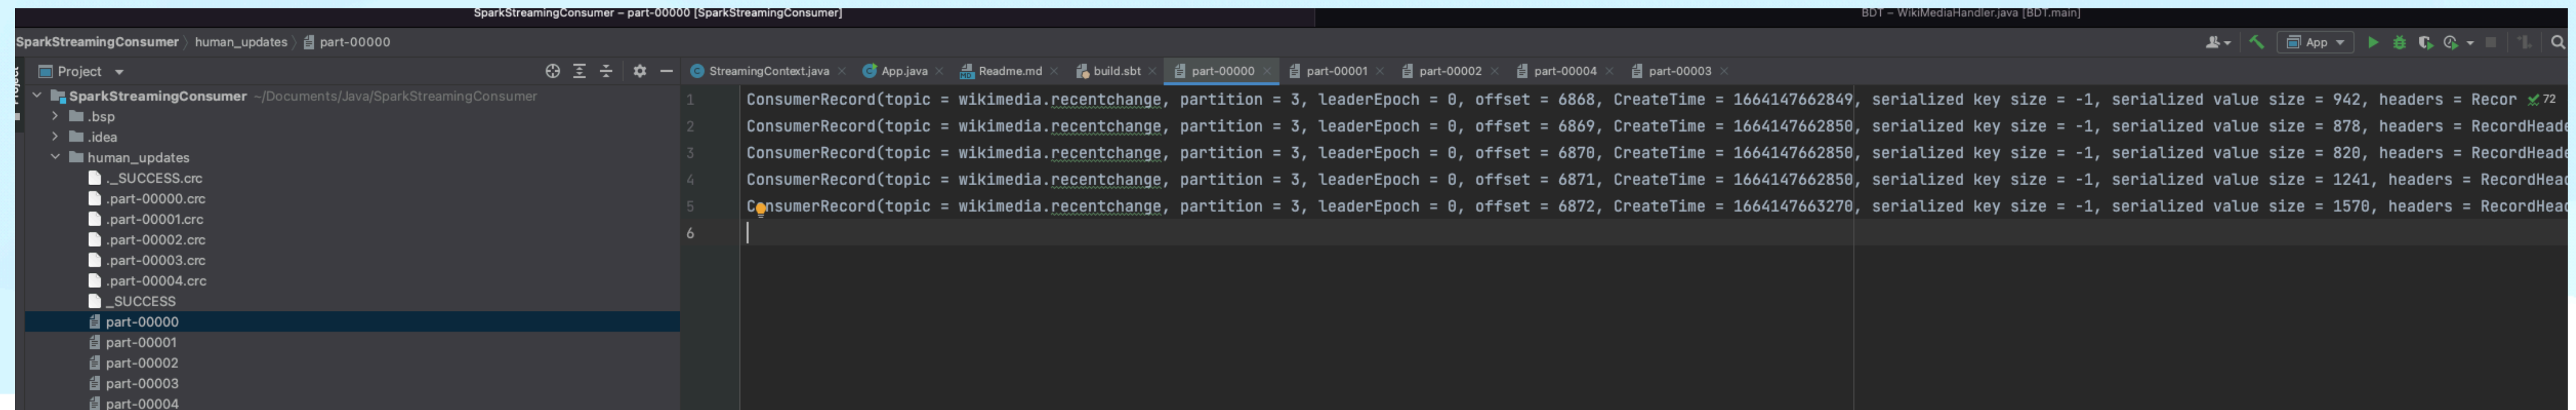
Spark RDD filtering and saving to HDFS

- Enable the stream and filter the output and save the final RDD result as text file.

```
public static void main(String[] args) throws InterruptedException {  
  
    JavaInputDStream<ConsumerRecord<String, String>> javaInputDStream = StreamingContext.javaInputDStream();  
    JavaDStream<ConsumerRecord<String, String>> human = javaInputDStream.filter(StreamingContext::isHuman);  
    javaInputDStream.foreachRDD(rdd -> {  
        //rdd.saveAsTextFile("human_updates");  
        rdd.saveAsTextFile("hdfs://cloudera:8020/user/cloudera/BDTPProject/human_updated");  
    });  
    StreamingContext.getOrCreateStreamContext(StreamingContext.getOrCreateStreamContext(),  
    StreamingContext.getJavaStreamingContext().awaitTermination();  
}
```

Result

The final text records in file system.



The screenshot displays an IDE interface. On the left, a project explorer shows a directory structure for 'SparkStreamingConsumer' with a subdirectory 'human_updates' containing several '.part-00000.crc' files. The main editor area shows a Java file named 'part-00000.crc' with five lines of code, each representing a 'ConsumerRecord' object. The records have a topic of 'wikimedia.recentchange', partition 3, and various offsets and timestamps. The IDE also shows other open files like 'StreamingContext.java', 'App.java', and 'Readme.md'.

```
1 ConsumerRecord(topic = wikimedia.recentchange, partition = 3, leaderEpoch = 0, offset = 6868, CreateTime = 1664147662849, serialized key size = -1, serialized value size = 942, headers = Recor ✓ 72
2 ConsumerRecord(topic = wikimedia.recentchange, partition = 3, leaderEpoch = 0, offset = 6869, CreateTime = 1664147662850, serialized key size = -1, serialized value size = 878, headers = RecordHead
3 ConsumerRecord(topic = wikimedia.recentchange, partition = 3, leaderEpoch = 0, offset = 6870, CreateTime = 1664147662850, serialized key size = -1, serialized value size = 820, headers = RecordHead
4 ConsumerRecord(topic = wikimedia.recentchange, partition = 3, leaderEpoch = 0, offset = 6871, CreateTime = 1664147662850, serialized key size = -1, serialized value size = 1241, headers = RecordHead
5 ConsumerRecord(topic = wikimedia.recentchange, partition = 3, leaderEpoch = 0, offset = 6872, CreateTime = 1664147663270, serialized key size = -1, serialized value size = 1570, headers = RecordHead
6 |
```